
Part 2: Defining the Machine Learning Problem (car price prediction dataset)

This dataset is a tabular dataset obtained from Kaggle in CSV format and contains information about used cars, such as manufacturing year, selling price, present market price, kilometers driven, fuel type, seller type, transmission, and ownership history. The dataset includes both numerical and categorical features, making it suitable for machine learning applications.

The objective of this project is to predict the selling price of a used car, which makes it a **regression problem**. The target variable is **Selling_Price**, and the model learns the relationship between the car's features and its selling price. This task uses **supervised machine learning** to accurately estimate car prices based on historical data.